

=====

MONGO ARCHIVER – Spring Boot (v3)

Configurable cross-collection MongoDB archiver with cron scheduling
and automatic temporary index management

=====

=====

PROJECT STRUCTURE

=====

```
mongo-archiver-java-v3/
├─ pom.xml
├─ README.md
├─ src/main/
│   └─ resources/
│       └─ application.yml           ← Only file you edit for new
entities
│   └─ java/com/archiver/
│       ├── MongoArchiverApplication.java    ← Entry point
│       ├── config/
│       │   └─ ArchiverProperties.java        ← @ConfigurationProperties
binding
│   ├── core/
│   │   └─ Archiver.java                     ← Core archive engine
│   ├── job/
│   │   └─ ArchiveCronJob.java               ← @Scheduled cron trigger
│   ├── model/
│   │   └─ ArchiveLogEntry.java              ← archiveLog document model
│   └─ util/
│       ├── ArchiveLogger.java               ← Writes outcomes to
archiveLog
│       ├── IdPathResolver.java              ← Dot-notation query builder
│       └─ IndexManager.java                 ← Temporary index lifecycle
(NEW)
```

=====

HOW IT WORKS

=====

1. ANCHOR QUERY

The anchor collection is queried for records where `createdTsField < now - archiveAfterDays`.

Only the anchor uses time as the archival gate. Related collections are never filtered by time – they are looked up purely by entity ID.

2. TEMPORARY INDEX MANAGEMENT (IndexManager)

Before any batching begins, IndexManager ensures an index exists on every field that will be queried:

- Anchor collection : createdTsField (cutoff query)
- Anchor collection : idField (per-entity archive query)
- Each related coll : idPath (per-entity query)

For each field:

- └ Index already exists? → Reuse it. Record as PRE-EXISTING. Never drop.
- └ No index exists? → Create one named "_arch_tmp_<field>".
Record as OWNED. Drop in finally block.

Cleanup is in a finally block in Archiver.run() – temporary indexes are ALWAYS dropped even if the run crashes mid-way.

3. BATCH PROCESSING

Matched entity IDs are split into chunks of batchSize to cap memory usage.

4. PER ENTITY – related collections first, anchor last

For each entity ID:

- a. All related collections are archived (find → bulk insert → optionally delete).
- b. Only if EVERY related collection succeeded → anchor is archived and deleted.
- c. If any related collection failed → anchor is LEFT UNTOUCHED.
The next createdTs sweep will naturally retry this entity.

5. DUPLICATE SAFETY

BulkOperations.UNORDERED is used for inserts. Duplicate _id errors on re-runs are caught and logged – the run is not aborted.

6. AUDIT LOG

Every outcome (SUCCESS / FAILED / NO_DOCUMENTS) is upserted into the archiveLog collection keyed on (jobName, entityId, collection).

TEMPORARY INDEX NAMING

Indexes created by this tool are named with the prefix: _arch_tmp_

Example names:

```

_arch_tmp_createdAt      on anchor.createdAt
_arch_tmp_orderId        on anchor.orderId
_arch_tmp_order_id       on orderPayments.order.id  (dots → underscores)
_arch_tmp_meta_referenceId on orderEvents.meta.referenceId

```

If a run crashes before cleanup, you can manually drop leftover indexes with:

```

db.<collection>.getIndexes()
db.<collection>.dropIndex("_arch_tmp_<fieldname>")

```

CONFIGURATION REFERENCE

Environment variables:

Variable	Default	Purpose
MONGO_URI	mongodb://localhost:27017	MongoDB URI
MONGO_DB	appdb	Database name

Common cron expressions (Spring 6-field format):

```

0 0 2 * * ?      Every day at 2:00 AM
0 0 */6 * * ?    Every 6 hours
0 0 2 ? * MON     Every Monday at 2:00 AM
0 0 2 1 * ?      First of every month at 2:00 AM

```

ID path examples:

```

"orderId"        matches { "orderId": "abc" }
"order.id"       matches { "order": { "id": "abc" } }
"meta.referenceId" matches { "meta": { "referenceId": "abc" } }
"payload.orderId" matches { "payload": { "orderId": "abc" } }

```

Adding a new entity type – edit application.yml only:

```

- name: invoices
  batch-size: 50
  anchor:
    collection: invoices
    id-field: invoiceId
    created-ts-field: issuedAt
    archive-after-days: 180
    delete-after-archive: true
    archive-collection: invoices_archive
  related:
    - collection: invoiceLineItems
      id-path: invoiceId

```

```
delete-after-archive: true
archive-collection: invoiceLineItems_archive
```

```
=====
RUNNING
=====
```

```
# Build
mvn clean package -DskipTests

# Scheduled mode
java -jar target/mongo-archiver-1.0.0.jar

# One-shot / backfill / testing
java -jar target/mongo-archiver-1.0.0.jar --run-now

# With custom connection
MONGO_URI=mongodb://prod-host:27017 MONGO_DB=mydb java -jar target/mongo-archiver-1.0.0.jar
```

```
=====
MONITORING
=====
```

```
Query the archiveLog collection:
```

```
// All failures
db.archiveLog.find({ status: "FAILED" }).sort({ runAt: -1 })
```

```
// Failures for a specific job
db.archiveLog.find({ jobName: "orders", status: "FAILED" })
```

```
// Full history for one entity
db.archiveLog.find({ entityId: "abc123" }).sort({ runAt: -1 })
```

```
// Summary by job + status
db.archiveLog.aggregate([
  { $group: { _id: { job: "$jobName", status: "$status" }, count: { $sum: 1 } } }
])
```

```
// Check for any leftover temporary indexes after a crash
db.orders.getIndexes().filter(i => i.name.startsWith("_arch_tmp_"))
```

```
=====
FILE: pom.xml
=====
```

```
<?xml version="1.0" encoding="UTF-8"?>
<project xmlns="http://maven.apache.org/POM/4.0.0"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://maven.apache.org/POM/4.0.0
    https://maven.apache.org/xsd/maven-4.0.0.xsd">
  <modelVersion>4.0.0</modelVersion>

  <parent>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-parent</artifactId>
    <version>3.2.3</version>
    <relativePath/>
  </parent>

  <groupId>com.archiver</groupId>
  <artifactId>mongo-archiver</artifactId>
  <version>1.0.0</version>
  <name>mongo-archiver</name>
  <description>Configurable cross-collection MongoDB archiver</description>

  <properties>
    <java.version>17</java.version>
  </properties>

  <dependencies>
    <!-- Spring Boot core (scheduling, config binding, logging) -->
    <dependency>
      <groupId>org.springframework.boot</groupId>
      <artifactId>spring-boot-starter</artifactId>
    </dependency>

    <!-- Spring Data MongoDB - provides MongoTemplate + MongoDBFactory -
->
    <dependency>
      <groupId>org.springframework.boot</groupId>
      <artifactId>spring-boot-starter-data-mongodb</artifactId>
    </dependency>

    <!-- Lombok - reduces boilerplate on config model classes -->
    <dependency>
      <groupId>org.projectlombok</groupId>
      <artifactId>lombok</artifactId>
      <optional>true</optional>
```

```

</dependency>

<!-- Generates IDE metadata for @ConfigurationProperties classes -->
<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-configuration-processor</artifactId>
    <optional>true</optional>
</dependency>

<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-test</artifactId>
    <scope>test</scope>
</dependency>
</dependencies>

<build>
    <plugins>
        <plugin>
            <groupId>org.springframework.boot</groupId>
            <artifactId>spring-boot-maven-plugin</artifactId>
            <configuration>
                <excludes>
                    <exclude>
                        <groupId>org.projectlombok</groupId>
                        <artifactId>lombok</artifactId>
                    </exclude>
                </excludes>
            </configuration>
        </plugin>
    </plugins>
</build>
</project>

```

```

=====
FILE: src/main/resources/application.yml
=====

```

```

spring:
  data:
    mongodb:
      uri: ${MONGO_URI:mongodb://localhost:27017}
      database: ${MONGO_DB:appdb}

archiver:

```

```

# Spring cron expression (6 fields: second minute hour day month weekday)
# Default: every day at 2:00 AM
schedule: "0 0 2 * * ?"
timezone: UTC

# Collection where archive run results are persisted for auditing
archive-log-collection: archiveLog

# — Archive jobs —————
# Each job = one logical entity type (e.g. orders, invoices, sessions).
# Add new entity types here – no Java changes needed.
jobs:

  - name: orders
    batch-size: 100

    # The anchor drives WHAT to archive and WHEN (via createdTs age).
    anchor:
      collection: orders
      id-field: orderId          # Top-level field on the anchor collection
      created-ts-field: createdAt # Timestamp used to decide if record is old
enough
      archive-after-days: 90
      delete-after-archive: true
      archive-collection: orders_archive

    # All other collections that share the same entity ID.
    # Each can declare a different field name/path for the ID.
    related:
      - collection: orderItems
        id-path: orderId          # Top-level
        delete-after-archive: true
        archive-collection: orderItems_archive

      - collection: orderPayments
        id-path: order.id          # Nested – dot notation supported
        delete-after-archive: true
        archive-collection: orderPayments_archive

      - collection: orderEvents
        id-path: meta.referenceId  # Deeply nested
        delete-after-archive: false # Copy only – keep source
        archive-collection: orderEvents_archive

      - collection: orderNotifications
        id-path: payload.orderId
        delete-after-archive: true

```

```

        archive-collection: orderNotifications_archive

# — Add more entity types below
# - name: invoices
#   batch-size: 50
#   anchor:
#     collection: invoices
#     id-field: invoiceId
#     created-ts-field: issuedAt
#     archive-after-days: 180
#     delete-after-archive: true
#     archive-collection: invoices_archive
#   related:
#     - collection: invoiceLineItems
#       id-path: invoiceId
#       delete-after-archive: true
#       archive-collection: invoiceLineItems_archive

logging:
  level:
    com.archiver: INFO
  pattern:
    console: "%d{yyyy-MM-dd HH:mm:ss} [%-5level] %logger{36} - %msg%n"
    file: "%d{yyyy-MM-dd HH:mm:ss} [%-5level] %logger{36} - %msg%n"
  file:
    name: logs/archiver.log
    max-size: 10MB
    max-history: 5

=====
FILE: src/main/java/com/archiver/MongoArchiverApplication.java
=====

package com.archiver;

import com.archiver.job.ArchiveCronJob;
import lombok.RequiredArgsConstructor;
import lombok.extern.slf4j.Slf4j;
import org.springframework.boot.ApplicationArguments;
import org.springframework.boot.ApplicationRunner;
import org.springframework.boot.SpringApplication;
import org.springframework.boot.autoconfigure.SpringBootApplication;
import org.springframework.boot.context.properties.EnableConfigurationProperties;
import org.springframework.scheduling.annotation.EnableScheduling;

/**

```



```

* MongoArchiverApplication
*
* <p>Entry point for the MongoDB archive service.
*
* <p>Modes:
* <ul>
*   <li><b>Scheduled (default)</b> – runs on the cron expression in
application.yml.</li>
*   <li><b>One-shot</b> – pass {@code --run-now} as a program argument to execute
*     immediately and exit. Useful for backfills or CI testing.</li>
* </ul>
*
* <p>Run:
* <pre>
*   # Scheduled mode
*   java -jar mongo-archiver.jar
*
*   # One-shot / backfill
*   java -jar mongo-archiver.jar --run-now
*
*   # With overridden connection strings
*   MONGO_PRIMARY_URI=mongodb://prod-host:27017 java -jar mongo-archiver.jar
* </pre>
*/
@Slf4j
@SpringBootApplication
@EnableScheduling
@RequiredArgsConstructor
public class MongoArchiverApplication implements ApplicationRunner {

    private final ArchiveCronJob archiveCronJob;

    public static void main(String[] args) {
        SpringApplication.run(MongoArchiverApplication.class, args);
    }

    /**
     * After full context startup: if {@code --run-now} was passed, execute
immediately and exit.
     */
    @Override
    public void run(ApplicationArguments args) {
        if (args.containsOption("run-now")) {
            log.info("[MongoArchiverApplication] --run-now flag detected;
executing immediately");
            archiveCronJob.runNow();
            log.info("[MongoArchiverApplication] One-shot run complete – shutting

```

```

down");
        System.exit(0);
    }

    log.info("[MongoArchiverApplication] Running in scheduled mode");
}
}

```

```

=====
FILE: src/main/java/com/archiver/config/ArchiverProperties.java
=====

```

```

package com.archiver.config;

import lombok.Data;
import org.springframework.boot.context.properties.ConfigurationProperties;
import org.springframework.stereotype.Component;

import java.util.List;

/**
 * Binds the entire {@code archiver.*} block from application.yml.
 *
 * Single-DB version – no connection registry, no per-collection connection names.
 * Spring Boot auto-configures MongoTemplate from spring.data.mongodb.*
 */
properties.
*/
@Data
@Component
@ConfigurationProperties(prefix = "archiver")
public class ArchiverProperties {

    /** Spring cron expression (6 fields). Default: 2 AM daily. */
    private String schedule = "0 0 2 * * ?";

    /** Timezone for the cron schedule. */
    private String timezone = "UTC";

    /** Collection name for archive run logs. */
    private String archiveLogCollection = "archiveLog";

    /** One entry per logical entity type to archive (orders, invoices, etc.). */
    private List<JobProperties> jobs;

    // — Nested config classes —————

```

```

@Data
public static class JobProperties {
    /** Human-readable name used in logs and archiveLog entries. */
    private String name;

    /** Number of entity IDs to process per cron tick. */
    private int batchSize = 100;

    /** The collection that drives the "old enough to archive?" decision. */
    private AnchorProperties anchor;

    /** All other collections that share the same entity ID. */
    private List<RelatedProperties> related;
}

@Data
public static class AnchorProperties {
    private String collection;

    /** Top-level field holding the entity ID in the anchor collection. */
    private String idField;

    /** Timestamp field used to measure record age. */
    private String createdTsField;

    /** Archive records older than this many days. */
    private int archiveAfterDays;

    /** Remove the source document after a successful archive write. */
    private boolean deleteAfterArchive = true;

    /** Name of the archive collection to write into (same DB). */
    private String archiveCollection;
}

@Data
public static class RelatedProperties {
    private String collection;

    /**
     * Dot-notation path to the entity ID within this collection's documents.
     * Examples: {@code "orderId"}, {@code "order.id"}, {@code
"meta.referenceId"}
     */
    private String idPath;

    private boolean deleteAfterArchive = true;
}

```

```

        /** Name of the archive collection to write into (same DB). */
        private String archiveCollection;
    }
}

```

```

=====
FILE: src/main/java/com/archiver/core/Archiver.java
=====

```

```

package com.archiver.core;

import com.archiver.config.ArchiverProperties;
import com.archiver.config.ArchiverProperties.AnchorProperties;
import com.archiver.config.ArchiverProperties.JobProperties;
import com.archiver.config.ArchiverProperties.RelatedProperties;
import com.archiver.model.ArchiveLogEntry;
import com.archiver.model.ArchiveLogEntry.Status;
import com.archiver.util.ArchiveLogger;
import com.archiver.util.IdPathResolver;
import com.archiver.util.IndexManager;
import lombok.RequiredArgsConstructor;
import lombok.extern.slf4j.Slf4j;
import org.bson.Document;
import org.springframework.data.mongodb.core.BulkOperations;
import org.springframework.data.mongodb.core.MongoTemplate;
import org.springframework.data.mongodb.core.query.Query;

import java.util.ArrayList;
import java.util.Calendar;
import java.util.Date;
import java.util.LinkedHashSet;
import java.util.List;
import java.util.Set;

/**
 * Archiver
 *
 * <p>Core archive engine for one configured job.
 *
 * <h3>Index strategy</h3>
 * <p>Before batching begins, an {@link IndexManager} ensures an index exists on
 * every queried field:
 * <ul>
 * <li>Anchor collection: {@code createdTsField} (cutoff query) and {@code
idField}

```

```

*      (per-entity archive query).</li>
*      <li>Each related collection: {@code idPath} (per-entity query).</li>
* </ul>
* If an index already exists on a field it is reused and never dropped.
* If a temporary index was created by this run, it is dropped in a {@code
finally}
* block – guaranteeing cleanup even on mid-run failures.
*
* <h3>Anchor-last strategy</h3>
* <p>Related collections are archived first. The anchor is archived and deleted
* only if every related collection succeeded. On any related failure the anchor
* stays in the source DB so the next {@code createdTs} sweep retries naturally.
*/
@Slf4j
@RequiredArgsConstructor
public class Archiver {

    private final JobProperties jobConfig;
    private final MongoTemplate mongoTemplate;
    private final ArchiveLogger archiveLogger;

    // — Entry point —————

    public void run() {
        String jobName = jobConfig.getName();
        Date cutoff = getCutoffDate(jobConfig.getAnchor().getArchiveAfterDays());

        log.info("[Archiver:{}] Starting – archiving records with {} < {}",
            jobName, jobConfig.getAnchor().getCreatedTsField(), cutoff);

        // One IndexManager per run – tracks only indexes created by THIS run
        IndexManager indexManager = new IndexManager(mongoTemplate);

        try {
            // — Step 1: ensure indexes on all queried fields before any batching
            ensureIndexes(indexManager);

            // — Step 2: fetch entity IDs from anchor —————
            List<String> ids = fetchAnchorIds(cutoff);
            if (ids.isEmpty()) {
                log.info("[Archiver:{}] No records to archive", jobName);
                return;
            }

            log.info("[Archiver:{}] {} entities to process", jobName, ids.size());

```

```

// — Step 3: process in batches —————
int batchSize = jobConfig.getBatchSize();
int totalBatches = (int) Math.ceil((double) ids.size() / batchSize);

for (int i = 0; i < ids.size(); i += batchSize) {
    List<String> batch = ids.subList(i, Math.min(i + batchSize,
ids.size()));
    log.info("[Archiver:{}] Batch {}/{} ({} IDs)",
        jobName, (i / batchSize) + 1, totalBatches, batch.size());
    processBatch(batch);
}

log.info("[Archiver:{}] Archive run complete", jobName);

} finally {
    // — Always drop temporary indexes — even on exception / partial
failure —
    indexManager.dropCreatedIndexes();
}
}

// — Index management —————

/**
 * Ensure an index exists on every field this job will query.
 * Pre-existing indexes are noted but never dropped.
 * Newly created indexes are owned by this run and dropped in the finally
block.
 */
private void ensureIndexes(IndexManager indexManager) {
    AnchorProperties anchor = jobConfig.getAnchor();

    // Anchor — two fields: the timestamp gate and the ID lookup
    indexManager.ensureIndex(anchor.getCollection(),
anchor.getCreatedTsField());
    indexManager.ensureIndex(anchor.getCollection(), anchor.getIdField());

    // Each related collection — the ID path used for per-entity queries
    for (RelatedProperties rel : jobConfig.getRelated()) {
        indexManager.ensureIndex(rel.getCollection(), rel.getIdPath());
    }

    log.info("[Archiver:{}] Index check complete — {} temporary index(es)
created",
        jobConfig.getName(), indexManager.ownedCount());
}

```

// — Batch processing —————

```
private void processBatch(List<String> entityIds) {
    for (String entityId : entityIds) {
        boolean allRelatedSucceeded = archiveRelatedCollections(entityId);

        if (allRelatedSucceeded) {
            archiveAnchor(entityId);
        } else {
            log.warn("[Archiver:{}] Skipping anchor for entity '{}' - " +
                "related collection failure. Will retry on next createdTs
sweep.",
                jobConfig.getName(), entityId);
        }
    }
}
```

```
private boolean archiveRelatedCollections(String entityId) {
    boolean allSucceeded = true;
    for (RelatedProperties rel : jobConfig.getRelated()) {
        boolean ok = archiveCollection(
            entityId,
            rel.getCollection(),
            rel.getIdPath(),
            rel.getArchiveCollection(),
            rel.isDeleteAfterArchive()
        );
        if (!ok) allSucceeded = false;
    }
    return allSucceeded;
}
```

```
private void archiveAnchor(String entityId) {
    AnchorProperties anchor = jobConfig.getAnchor();
    archiveCollection(
        entityId,
        anchor.getCollection(),
        anchor.getIdField(),
        anchor.getArchiveCollection(),
        anchor.isDeleteAfterArchive()
    );
}
```

// — Core archive operation —————

/**

* Find all documents for {@code entityId} in {@code srcCollection},

```

    * write them to {@code archiveCollection}, optionally delete from source.
    *
    * @return true on success or no-documents; false on error
    */
private boolean archiveCollection(
    String entityId,
    String srcCollection,
    String idPath,
    String archiveCollection,
    boolean deleteAfterArchive
) {
    String jobName = jobConfig.getName();
    try {
        Query query = IdPathResolver.buildSingleQuery(idPath, entityId);
        List<Document> docs = mongoTemplate.find(query, Document.class,
srcCollection);

        if (docs.isEmpty()) {
            log.debug("[Archiver:{}] No documents in '{}' for entity '{}'",
                jobName, srcCollection, entityId);
            writeLog(jobName, entityId, srcCollection, Status.NO_DOCUMENTS, 0,
null);

            return true;
        }

        // Stamp with archive metadata
        Date archivedAt = new Date();
        docs.forEach(doc -> {
            doc.append("_archivedAt", archivedAt);
            doc.append("_archivedFrom", srcCollection);
        });

        // Bulk insert - UNORDERED tolerates duplicate _id on re-runs
        int inserted = bulkInsert(docs, archiveCollection, jobName, entityId,
srcCollection);

        if (deleteAfterArchive) {
            mongoTemplate.remove(query, srcCollection);
        }

        log.info("[Archiver:{}] Archived {} doc(s): '{}' → '{}' (entity:
'{}')",
            jobName, docs.size(), srcCollection, archiveCollection, entityId);

        writeLog(jobName, entityId, srcCollection, Status.SUCCESS,
docs.size(), null);
        return true;
    }
}

```



```

    } catch (Exception e) {
        log.error("[Archiver:{}] FAILED '{}' for entity '{}': {}",
            jobName, srcCollection, entityId, e.getMessage(), e);
        writeLog(jobName, entityId, srcCollection, Status.FAILED, 0,
e.getMessage());
        return false;
    }
}

// — Bulk insert with duplicate tolerance —————

private int bulkInsert(
    List<Document> docs, String collection,
    String jobName, String entityId, String srcCollection
) {
    try {
        BulkOperations bulk =
mongoTemplate.bulkOps(BulkOperations.BulkMode.UNORDERED, collection);
        docs.forEach(bulk::insert);
        return bulk.execute().getRawResult().getInsertedCount();

        } catch (org.springframework.data.mongodb.BulkOperationException e) {
            int inserted = e.getResult() != null ?
e.getResult().getInsertedCount() : 0;
            int dupes = docs.size() - inserted;
            if (dupes > 0) {
                log.warn("[Archiver:{}] '{}' entity='{}' - {} already archived, {}
newly inserted",
                    jobName, srcCollection, entityId, dupes, inserted);
            }
            return inserted;
        }
    }

// — Anchor ID fetching —————

private List<String> fetchAnchorIds(Date cutoff) {
    AnchorProperties anchor = jobConfig.getAnchor();
    try {
        Query query =
IdPathResolver.buildCutoffQuery(anchor.getCreatedTsField(), cutoff);
        query.fields().include(anchor.getIdField()).exclude("_id");

        Set<String> ids = new LinkedHashSet<>();
        mongoTemplate.find(query, Document.class, anchor.getCollection())
            .forEach(doc -> {

```

```

        String id = IdPathResolver.extractIdAsString(doc,
anchor.getIdField());
        if (id != null) ids.add(id);
    });

    log.info("[Archiver:{}] Found {} anchor IDs in '{}'",
        jobConfig.getName(), ids.size(), anchor.getCollection());
    return new ArrayList<>(ids);

} catch (Exception e) {
    log.error("[Archiver:{}] Failed to fetch anchor IDs from '{}': {}",
        jobConfig.getName(), anchor.getCollection(), e.getMessage(), e);
    return List.of();
}
}

// ——— Helpers —————

private Date getCutoffDate(int days) {
    Calendar cal = Calendar.getInstance();
    cal.add(Calendar.DATE, -days);
    return cal.getTime();
}

private void writeLog(String jobName, String entityId, String collection,
    Status status, int count, String error) {
    archiveLogger.log(ArchiveLogEntry.builder()
        .jobName(jobName)
        .entityId(entityId)
        .collection(collection)
        .status(status)
        .archivedCount(count)
        .errorMessage(error)
        .runAt(new Date())
        .build());
}
}

```

```

=====
FILE: src/main/java/com/archiver/job/ArchiveCronJob.java
=====

```

```

package com.archiver.job;

import com.archiver.config.ArchiverProperties;
import com.archiver.core.Archiver;

```

```

import com.archiver.util.ArchiveLogger;
import jakarta.annotation.PostConstruct;
import lombok.RequiredArgsConstructor;
import lombok.extern.slf4j.Slf4j;
import org.springframework.data.mongodb.core.MongoTemplate;
import org.springframework.scheduling.annotation.Scheduled;
import org.springframework.stereotype.Component;

import java.time.Instant;
import java.util.concurrent.atomic.AtomicBoolean;

/**
 * ArchiveCronJob
 *
 * Triggers the archive run on schedule. Delegates to one {@link Archiver}
 * instance per configured job. All jobs share the same auto-configured
 * {@link MongoTemplate} (single DB).
 */
@Slf4j
@Component
@RequiredArgsConstructor
public class ArchiveCronJob {

    private final ArchiverProperties properties;
    private final MongoTemplate mongoTemplate;
    private final ArchiveLogger archiveLogger;

    /** Prevents overlapping runs if a job takes longer than the cron interval. */
    private final AtomicBoolean running = new AtomicBoolean(false);

    @PostConstruct
    public void init() {
        archiveLogger.ensureIndexes();
        log.info("[ArchiveCronJob] Initialised. Schedule='{}' timezone='{}'",
            properties.getSchedule(), properties.getTimezone());
        log.info("[ArchiveCronJob] {} job(s) configured: {}",
            properties.getJobs().size(),

properties.getJobs().stream().map(ArchiverProperties.JobProperties::getName).toList
        }

        @Scheduled(cron = "${archiver.schedule}", zone = "${archiver.timezone}")
        public void runScheduled() {
            execute("scheduled");
        }

        /** Programmatic trigger – useful for one-shot runs or testing. */

```

```

public void runNow() {
    execute("manual");
}

// — Private —————

private void execute(String trigger) {
    if (!running.compareAndSet(false, true)) {
        log.warn("[ArchiveCronJob] Previous run still in progress – skipping
{} trigger", trigger);
        return;
    }

    long start = System.currentTimeMillis();
    log.info("[ArchiveCronJob] Run started at {} (trigger: {})",
Instant.now(), trigger);

    try {
        for (ArchiverProperties.JobProperties jobConfig :
properties.getJobs()) {
            try {
                new Archiver(jobConfig, mongoTemplate, archiveLogger).run();
            } catch (Exception e) {
                log.error("[ArchiveCronJob] Unhandled error in job '{}': {}",
jobConfig.getName(), e.getMessage(), e);
            }
        }
    } finally {
        running.set(false);
        log.info("[ArchiveCronJob] Completed in {}s",
String.format("%.1f", (System.currentTimeMillis() - start) /
1000.0));
    }
}
}

```

```

=====
FILE: src/main/java/com/archiver/model/ArchiveLogEntry.java
=====

```

```

package com.archiver.model;

import lombok.Builder;
import lombok.Data;

import java.util.Date;

```

```

/**
 * One row in the {@code archiveLog} collection.
 * Keyed on (jobName, entityId, collection) – upserted on every run.
 *
 * {
 *   jobName      : "orders",
 *   entityId     : "abc123",
 *   collection   : "orderPayments",
 *   status       : "SUCCESS" | "FAILED" | "NO_DOCUMENTS",
 *   archivedCount: 5,
 *   errorMessage : null,
 *   runAt        : ISODate("..."),
 *   retryCount   : 2
 * }
 */
@Data
@Builder
public class ArchiveLogEntry {

    public enum Status {
        SUCCESS,
        FAILED,
        NO_DOCUMENTS
    }

    private String jobName;
    private String entityId;
    private String collection;
    private Status status;
    private int archivedCount;
    private String errorMessage;
    private Date runAt;
}

```

```

=====
FILE: src/main/java/com/archiver/util/ArchiveLogger.java
=====

```

```

package com.archiver.util;

import com.archiver.config.ArchiverProperties;
import com.archiver.model.ArchiveLogEntry;
import lombok.RequiredArgsConstructor;
import lombok.extern.slf4j.Slf4j;
import org.bson.Document;

```

```

import org.springframework.data.domain.Sort;
import org.springframework.data.mongodb.core.MongoTemplate;
import org.springframework.data.mongodb.core.index.CompoundIndexDefinition;
import org.springframework.data.mongodb.core.index.Index;
import org.springframework.data.mongodb.core.query.Criteria;
import org.springframework.data.mongodb.core.query.Query;
import org.springframework.data.mongodb.core.query.Update;
import org.springframework.stereotype.Component;

import java.util.Date;

/**
 * ArchiveLogger
 *
 * Persists the outcome of each archive operation to the {@code archiveLog}
 * collection for auditing. Uses Spring Boot's auto-configured {@link
MongoTemplate}.
 *
 * Upserts keyed on (jobName, entityId, collection) so re-runs overwrite cleanly.
 */
@Slf4j
@Component
@RequiredArgsConstructor
public class ArchiveLogger {

    private final MongoTemplate mongoTemplate;
    private final ArchiverProperties properties;

    /**
     * Upsert a log entry for one (entity, collection) outcome.
     * Silently swallows exceptions so logging never crashes the archive job.
     */
    public void log(ArchiveLogEntry entry) {
        try {
            Query query = Query.query(
                Criteria.where("jobName").is(entry.getJobName())
                    .and("entityId").is(entry.getEntityId())
                    .and("collection").is(entry.getCollection())
            );

            Update update = new Update()
                .set("status", entry.getStatus().name())
                .set("archivedCount", entry.getArchivedCount())
                .set("errorMessage", entry.getErrorMessage())
                .set("runAt", new Date())
                .inc("retryCount", 1)
                .setOnInsert("createdAt", new Date());

```

```

        mongoTemplate.upsert(query, update,
properties.getArchiveLogCollection());

        } catch (Exception e) {
            log.warn("[ArchiveLogger] Failed to write log for entity={}
collection={}: {}",
                entry.getEntityId(), entry.getCollection(), e.getMessage());
        }
    }

    /**
     * Ensure indexes on the archiveLog collection. Idempotent – safe to call on
     every startup.
     */
    public void ensureIndexes() {
        try {
            String col = properties.getArchiveLogCollection();

            mongoTemplate.indexOps(col).ensureIndex(
                new CompoundIndexDefinition(
                    new Document("jobName", 1).append("entityId",
1).append("collection", 1)
                ).unique()
            );
            mongoTemplate.indexOps(col).ensureIndex(
                new CompoundIndexDefinition(new Document("status",
1).append("jobName", 1))
            );
            mongoTemplate.indexOps(col).ensureIndex(
                new Index("runAt", Sort.Direction.DESC)
            );

            log.info("[ArchiveLogger] Indexes ensured on '{}'", col);
        } catch (Exception e) {
            log.warn("[ArchiveLogger] Index creation warning: {}",
e.getMessage());
        }
    }
}

```

```

=====
FILE: src/main/java/com/archiver/util/IdPathResolver.java
=====

```

```

package com.archiver.util;

```

```

import org.bson.Document;
import org.springframework.data.mongodb.core.query.Criteria;
import org.springframework.data.mongodb.core.query.Query;

import java.util.Date;
import java.util.List;

/**
 * IdPathResolver
 *
 * <p>Builds Spring Data MongoDB {@link Query} objects from dot-notation ID paths,
 * and walks in-memory {@link Document} objects to extract values at those paths.
 *
 * <p>MongoDB's query engine natively supports dot-notation in field names, so
 * {@code Criteria.where("meta.referenceId")} correctly matches nested documents.
 *
 * <p>Examples:
 * <pre>
 *     resolveValue(doc, "orderId")           → "abc123"
 *     resolveValue(doc, "order.id")          → "abc123"
 *     resolveValue(doc, "meta.referenceId") → "abc123"
 *
 *     buildSingleQuery("order.id", "abc")    → Query: { "order.id": "abc" }
 *     buildBatchQuery("order.id", ids)       → Query: { "order.id": { $in: [...] } }
 *     buildCutoffQuery("createdAt", date)    → Query: { "createdAt": { $lt: date } }
 * </pre>
 */
public final class IdPathResolver {

    private IdPathResolver() {}

    /**
     * Walk a dot-notation path on a {@link Document}, returning the nested value.
     *
     * @param doc source document
     * @param path dot-notation path, e.g. {@code "meta.referenceId"}
     * @return the resolved value, or {@code null} if any segment is missing
     */
    public static Object resolveValue(Document doc, String path) {
        String[] parts = path.split("\\.");
        Object current = doc;
        for (String part : parts) {
            if (!(current instanceof Document nested)) {
                return null;
            }
            current = nested.get(part);
        }
    }

```



```

        }
        return current;
    }

    /**
     * Build a Spring Data {@link Query} matching a single entity ID at the given
path.
     *
     * @param idPath    dot-notation field path, e.g. {@code "order.id"}
     * @param idValue    the ID value to match
     * @return {@link Query}
     */
    public static Query buildSingleQuery(String idPath, Object idValue) {
        return Query.query(Criteria.where(idPath).is(idValue));
    }

    /**
     * Build a Spring Data {@link Query} matching a batch of entity IDs via {@code
$in}.
     *
     * @param idPath    dot-notation field path, e.g. {@code "order.id"}
     * @param idValues    list of ID values
     * @return {@link Query} using {@code $in}
     */
    public static Query buildBatchQuery(String idPath, List<?> idValues) {
        return Query.query(Criteria.where(idPath).in(idValues));
    }

    /**
     * Build a Spring Data {@link Query} for records where a date field is before
     * the given cutoff. Used to query the anchor collection by {@code createdTs}.
     *
     * @param tsField    timestamp field name (top-level)
     * @param cutoff    upper bound (exclusive)
     * @return {@link Query}
     */
    public static Query buildCutoffQuery(String tsField, Date cutoff) {
        return Query.query(Criteria.where(tsField).lt(cutoff));
    }

    /**
     * Extract the value of a top-level field from a document as a String.
     * Returns {@code null} if the field is absent or null.
     *
     * @param doc    source document
     * @param field    top-level field name
     * @return string representation, or {@code null}

```

```

        */
    public static String extractIdAsString(Document doc, String field) {
        Object value = doc.get(field);
        return value != null ? value.toString() : null;
    }
}

```

```

=====
FILE: src/main/java/com/archiver/util/IndexManager.java
=====

```

```

package com.archiver.util;

import lombok.RequiredArgsConstructor;
import lombok.extern.slf4j.Slf4j;
import org.springframework.data.mongodb.core.MongoTemplate;
import org.springframework.data.mongodb.core.index.IndexInfo;
import org.springframework.data.mongodb.core.index.IndexOperations;
import
org.springframework.data.mongodb.core.index.MongoPersistentEntityIndexResolver;
import org.springframework.data.mongodb.core.index.PartialIndexFilter;
import org.springframework.data.mongodb.core.index.Index;

import java.util.LinkedHashSet;
import java.util.List;
import java.util.Set;

/**
 * IndexManager
 *
 * <p>Manages temporary indexes created for the duration of one archive run.
 *
 * <p>Behaviour per field:
 * <ul>
 *   <li>If an index on the field already exists → record it as <b>pre-
existing</b>,
 *       use it as-is, and <b>never drop it</b> after the run.</li>
 *   <li>If no index exists → create one, record it as <b>owned</b>,
 *       and drop it in {@link #dropCreatedIndexes()} when the run finishes.</li>
 * </ul>
 *
 * <p>{@link #dropCreatedIndexes()} must be called in a {@code finally} block so
 * that temporary indexes are always cleaned up even on mid-run failures.
 *
 * <p>One {@code IndexManager} instance per archive run – not a Spring singleton.
 */

```

```

@Slf4j
@RequiredArgsConstructor
public class IndexManager {

    /**
     * Prefix applied to every index name this manager creates.
     * Makes temporary indexes easy to identify if a run crashes before cleanup.
     */
    private static final String INDEX_PREFIX = "_arch_tmp_";

    private final MongoTemplate mongoTemplate;

    /**
     * Tracks indexes this manager created (as "collection::indexName" keys).
     * Pre-existing indexes are never added here and are never dropped.
     */
    private final Set<String> ownedIndexes = new LinkedHashSet<>();

    // — Public API —————

    /**
     * Ensure an index exists on {@code fieldPath} in {@code collection}.
     *
     * <p>If an index covering this field already exists (under any name), it is
     * reused and will not be dropped after the run. If no such index exists, a
     * temporary one is created and tracked for cleanup.
     *
     * @param collection MongoDB collection name
     * @param fieldPath dot-notation field path, e.g. {@code "meta.referenceId"}
     */
    public void ensureIndex(String collection, String fieldPath) {
        try {
            IndexOperations indexOps = mongoTemplate.indexOps(collection);

            if (indexAlreadyExists(indexOps, fieldPath)) {
                log.info("[IndexManager] Pre-existing index found on '{}.{}' -
will not drop after run",
                    collection, fieldPath);
                return;
            }

            // No existing index - create a temporary one
            String indexName = INDEX_PREFIX + fieldPath.replace(".", "_");
            indexOps.ensureIndex(
                new Index().on(fieldPath,
org.springframework.data.domain.Sort.Direction.ASC)
                    .named(indexName)

```

```

    );

    String key = collection + "::-" + indexName;
    ownedIndexes.add(key);
    log.info("[IndexManager] Created temporary index '{}' on '{}.{}'.'",
        indexName, collection, fieldPath);

    } catch (Exception e) {
        // Non-fatal - log and continue. Query will still work, just without
the index.
        log.warn("[IndexManager] Could not ensure index on '{}.{}': {}",
            collection, fieldPath, e.getMessage());
    }
}

/**
 * Drop all indexes that this manager created during the current run.
 * Pre-existing indexes are never touched.
 *
 * <p>Must be called in a {@code finally} block to guarantee cleanup on
failure.
 */
public void dropCreatedIndexes() {
    if (ownedIndexes.isEmpty()) {
        return;
    }

    log.info("[IndexManager] Dropping {} temporary index(es)...",
ownedIndexes.size());

    for (String key : ownedIndexes) {
        String[] parts = key.split("::-", 2);
        String collection = parts[0];
        String indexName = parts[1];
        try {
            mongoTemplate.indexOps(collection).dropIndex(indexName);
            log.info("[IndexManager] Dropped temporary index '{}' on '{}.{}'.'",
indexName, collection);
        } catch (Exception e) {
            log.warn("[IndexManager] Failed to drop index '{}' on '{}': {}",
                indexName, collection, e.getMessage());
        }
    }

    ownedIndexes.clear();
}

```

```

/**
 * How many temporary indexes this manager has created (useful for logging).
 */
public int ownedCount() {
    return ownedIndexes.size();
}

// — Private helpers —————

/**
 * Check whether any existing index on the collection covers {@code fieldPath}
 * as its leading (or only) key.
 *
 * <p>MongoDB represents dot-notation fields as-is in index key documents,
 * so {@code "meta.referenceId"} in the index info will match the path
 * exactly.
 */
private boolean indexAlreadyExists(IndexOperations indexOps, String fieldPath)
{
    try {
        List<IndexInfo> existing = indexOps.getIndexInfo();
        for (IndexInfo info : existing) {
            boolean covers = info.getIndexFields().stream()
                .anyMatch(f -> f.getKey().equals(fieldPath));
            if (covers) {
                return true;
            }
        }
    } catch (Exception e) {
        log.warn("[IndexManager] Could not read index info: {}",
e.getMessage());
    }
    return false;
}
}

```

```

=====
END OF FILE
=====

```